






















Empirical variance of ψ_{05} — Oracle vs Super-learner nuisances



Empirical median of ψ_{05} — Oracle vs Super-learner nuisances









Empirical MSE of ψ_{05} — Oracle vs Super-learner nuisances



Empirical coverage of ψ_{05} — Oracle vs Super-learner nuisances

Oracle
 Super-learner



Both runs are post-fix (2026-09-02): ee_three_step_ipw equation fix and the full-data covariate-plug-in working_model() change (mu05..mu45 fit on the full cross-fitting fold with A_t as a covariate, rather than only the $A_t == A_{t-1}$ subset) are both reflected here -- see CLAUDE.md.

Known residual gap: at $ph=0.3$, Simple g/Three-step-g's Super-learner variance still runs $\sim 3-4x$ above Oracle (Weighted regression's plain-parametric nuisances track Oracle closely everywhere). A matched $n=144,000$ check (40 iters, not yet at production 300-iter precision) did not show this gap closing with n .